

UCARE Research

Ray Andrew

November 5, 2019

Abstract

The abstract text goes here.

Contents

0.1	Todos	2
0.1.1	03 November 2019	2
0.2	Meeting Notes	3
0.2.0.1	09 October 2019	3
0.2.0.2	22 October 2019	3
0.2.0.3	30 October 2019	3
0.3	Summarization	4
0.3.1	JVM Hotspot	4
0.3.1.1	Oop Hierarchy	4
0.3.1.2	Oop Memory Layout	4
0.3.1.3	JVM Alignment	6
0.3.1.4	Array Allocation	8
0.3.1.5	Array Copy Mechanism	8
0.3.1.6	Type Array Inheritance	26
0.3.1.7	TLAB	27
0.3.1.7.1	Inheritance	27
0.3.1.7.2	Initialization	27
0.3.1.7.3	Notes	27
0.4	Notes	29
0.4.1	Tuesday, 08 October 2019	30
0.4.1.1	Kafka Bug Study	30
0.4.1.2	JVM Deduplication	30
0.4.2	Friday, 04 October 2019	31
0.4.2.1	Allocation Classes	31
0.4.2.2	Notes	31
0.4.3	Saturday, 19 October 2019	32
0.4.3.1	JVM Deduplication	32
0.5	Resources	34
0.5.1	Kernel Samepage Merging	34

0.1 Todos

0.1.1 03 November 2019

- Add markOop memory layout
- Upload docs into github
- Test dedup System.arraycopy

0.2 Meeting Notes

0.2.0.1 09 October 2019

- Try the Kafka sytem inside Eclipse OpenJ9 test whether parameter class-datashare is working or not
- IBM Java Multitenancy
- See implementation of native function (arrayCopy) and try see how the copy the array and look into the content
- No matplotlib and try to change it GNUPlot

0.2.0.2 22 October 2019

- Create our own SLABs, 32 Nodes, 7 Workloads
- Check if object is unmodified
- Understand what is Klass, MarkOop, Array looks like on memory layout
- send me string dedup paper, read it and we discuss it next tuesday

0.2.0.3 30 October 2019

- 1 more week to implement arraycopy to check whether we can call slab allocation directly (basically mmap)
- Mark class memory layout
- Upload the documents
- Check implementation of String Deduplication in OpenJ9
- Continuing other bugs

0.3 Summarization

0.3.1 JVM Hotspot

0.3.1.1 Oop Hierarchy

```
typedef class oopDesc*           oop;
typedef class instanceOopDesc*   instanceOop;
typedef class arrayOopDesc*      arrayOop;
typedef class objArrayOopDesc*   objArrayOop;
typedef class typeArrayOopDesc*  typeArrayOop;
```

0.3.1.2 Oop Memory Layout

1. narrowKlass = int, sizeof = 4
2. `const int HeapWordSize = 8;`
3. `const int LogBytesPerLong = 3;`
4. `const int BytesPerLong = 1 << LogBytesPerLong;`
5. `const int HeapWordsPerLong = BytesPerLong / HeapWordSize;`
6. Instance Oops for Object

Type	Name	64bit		Compressed	
		Offset	Size	Offset	Size
Header	mark *ptr	0	8	0	8
	klass *ptr	8	8	8	4
Fields	first field	16	n	12	n
	...fields	16 + n	...	12 + n	...

7. Instance Oops for Array

Type	Name	64bit		Compressed	
		Offset	Size	Offset	Size
Header	mark *ptr	0	8	0	8
	klass *ptr	8	8	8	4
Fields	length	16	4	12	4
	first field	24	n	16	n
	...fields	24 + n	...	16 + n	...

8. Klass

Data type	Name	64bit		Compressed	
		Offset	Size	Offset	Size
pointer	vtable *ptr	0	8	0	8
int	_layout_helper	8	4	8	4

9. Array Klass

Data type	Name	64bit		Compressed	
		Offset	Size	Offset	Size
pointer	vtable *ptr	0	8	0	8
int	_layout_helper	8	4	8	4

10. Mark Normal Object

Name	64bit		Name	Compressed	
	Offset	Size		Offset	Size
unused	0	25	unused	0	25
hash	25	31	hash	25	31
unused	56	1	unused	56	1
age	57	4	age	57	4
biased_lock	61	1	biased_lock	61	1
lock	62	2	lock	62	2

11. Mark Biased Object

Name	64bit		Name	Compressed	
	Offset	Size		Offset	Size
JavaThread*	0	54	JavaThread*	0	5
epoch	54	2	epoch	54	2
unused	56	1	unused	56	1
age	57	4	age	57	4
biased_lock	61	1	biased_lock	61	1
lock	62	2	lock	62	2

12. Mark CMS Promoted Object

Name	64bit		Name	Compressed	
	Offset	Size		Offset	Size
PromotedObject*	0	61	narrowOop	0	32
			unused	32	24
			cms_free	56	1
			unused	57	4
promo_bits	61	3	promo_bits	61	3

13. Mark CMS Free Object

Name	64bit		Name	Compressed	
	Offset	Size		Offset	Size
size	0	64	unused	0	21
			size	21	35
			cms_free	56	1
			unused	57	7

0.3.1.3 JVM Alignment

```
// Signed variants of alignment helpers.  There are two versions
↳ of each, a macro
// for use in places like enum definitions that require
↳ compile-time constant
// expressions and a function for all other places so as to get
↳ type checking.

// Using '(what) & ~align_mask(alignment)' to align 'what' down
↳ is broken when
// 'alignment' is an unsigned int and 'what' is a wider type. The
↳ & operation
// will widen the inverted mask, and not sign extend it, leading
↳ to a mask with
// zeros in the most significant bits. The use of
↳ align_mask_widened() solves
// this problem.
#define align_mask(alignment) ((alignment) - 1)
#define widen_to_type_of(what, type_carrier) (true ? (what) :
↳ (type_carrier))
#define align_mask_widened(alignment, type_carrier)
↳ widen_to_type_of(align_mask(alignment), (type_carrier))

#define align_down_(size, alignment) ((size) &
↳ ~align_mask_widened((alignment), (size)))
```



```

#define align_up_(size, alignment) (align_down_((size) +
↪ align_mask(alignment), (alignment)))

#define is_aligned_(size, alignment) (((size) &
↪ align_mask(alignment)) == 0)

// Temporary declaration until this file has been restructured.
template <typename T>
bool is_power_of_2_t(T x) {
    return (x != T(0)) && ((x & (x - 1)) == T(0));
}

// Helpers to align sizes and check for alignment

template <typename T, typename A>
inline T align_up(T size, A alignment) {
    assert(is_power_of_2_t(alignment), "must be a power of 2: "
↪ UINT64_FORMAT, (uint64_t)alignment);

    T ret = align_up_(size, alignment);
    assert(is_aligned_(ret, alignment), "must be aligned: "
↪ UINT64_FORMAT, (uint64_t)ret);

    return ret;
}

template <typename T, typename A>
inline T align_down(T size, A alignment) {
    assert(is_power_of_2_t(alignment), "must be a power of 2: "
↪ UINT64_FORMAT, (uint64_t)alignment);

    T ret = align_down_(size, alignment);
    assert(is_aligned_(ret, alignment), "must be aligned: "
↪ UINT64_FORMAT, (uint64_t)ret);

    return ret;
}


template <typename T, typename A>
inline bool is_aligned(T size, A alignment) {
    assert(is_power_of_2_t(alignment), "must be a power of 2: "
↪ UINT64_FORMAT, (uint64_t)alignment);

    return is_aligned_(size, alignment);
}

template <typename T>
inline T align_object_offset(T offset) {
    return align_up(offset, HeapWordsPerLong);
}

```

0.3.1.4 Array Allocation



1. `void OptoRuntime::new_array_C(Klass* array_type, int len, JavaThread* thread)`
2. `typeArrayOop oopFactory::new_typeArray(BasicType type, int length, TRAPS)`
3. `TypeArrayKlass* TypeArrayKlass::allocate(ClassLoaderData* loader_data, BasicType type, Symbol* name, TRAPS)`
4. `TypeArrayKlass* TypeArrayKlass::allocate(ClassLoaderData* loader_data, BasicType type, Symbol* name, TRAPS)`
5. `objArrayOop oopFactory::new_objArray(Klass* klass, int length, TRAPS)`

0.3.1.5 Array Copy Mechanism

- hotspot/share/prism/jvm.cpp

```
JVM_ENTRY(void, JVM_ArrayCopy(JNIEnv *env, jclass ignored,
↪ jobject src, jint src_pos, jobject dst, jint dst_pos, jint
↪ length))
    JVMWrapper("JVM_ArrayCopy");
    // Check if we have null pointers
    if (src == NULL || dst == NULL) {
        THROW(vmSymbols::java_lang_NullPointerException());
    }

    arrayOop s = arrayOop(JNIHandles::resolve_non_null(src));
    arrayOop d = arrayOop(JNIHandles::resolve_non_null(dst));
    assert(oopDesc::is_oop(s), "JVM_ArrayCopy: src not an oop");
    assert(oopDesc::is_oop(d), "JVM_ArrayCopy: dst not an oop");
    // Do copy
    s->klass()->copy_array(s, src_pos, d, dst_pos, length,
↪ thread);
JVM_END
```

- hotspot/share/oops/arrayOopDesc.hpp
based on the [subsubsection 0.3.1.1](#), array oopDesc can be both typeArrayOopDesc or objArrayOopDesc

```
class arrayOopDesc : public oopDesc { ... }

// so basically oopDesc have attribute method call klass()
Klass* oopDesc::klass() const {
    if (UseCompressedClassPointers) {
        return
            ↪ Klass::decode_class_not_null(_metadata._compressed_klass);
    } else {
        return _metadata._klass;
    }
}
```

- hotspot/share/oops/typeArrayKlass.hpp

```
void TypeArrayKlass::copy_array(arrayOop s, int src_pos,
    ↪ arrayOop d, int dst_pos, int length, TRAPS) {
    assert(s->is_typeArray(), "must be type array");

    // Check destination type.
    if (!d->is_typeArray()) {
        ResourceMark rm(THREAD);
        stringstream ss;
        if (d->is_objArray()) {
            ss.print("arraycopy: type mismatch: can not copy %s[]
                ↪ into object array[]",
                ↪ type2name_tab[ArrayKlass::cast(s->klass())->element_type()]);
        } else {
            ss.print("arraycopy: destination type %s is not an
                ↪ array", d->klass()->external_name());
        }

        THROW_MSG(vmSymbols::java_lang_ArrayStoreException(),
            ↪ ss.as_string());
    }

    if (element_type() !=
        ↪ TypeArrayKlass::cast(d->klass())->element_type()) {
        ResourceMark rm(THREAD);
        stringstream ss;
        ss.print("arraycopy: type mismatch: can not copy %s[] into
            ↪ %s[]",
            ↪ type2name_tab[ArrayKlass::cast(s->klass())->element_type()],
            ↪ type2name_tab[ArrayKlass::cast(d->klass())->element_type()]);
    }
```

```

        THROW_MSG(vmSymbols::java_lang_ArrayStoreException(),
        ↪ ss.as_string());
    }

    // Check if all offsets and lengths are non negative.
    if (src_pos < 0 || dst_pos < 0 || length < 0) {
        // Pass specific exception reason.
        ResourceMark rm(THREAD);
        stringstream ss;
        if (src_pos < 0) {
            ss.print("arraycopy: source index %d out of bounds for
            ↪ %s[%d]", src_pos,
            ↪ type2name_tab[ArrayKlass::cast(s->klass())->element_type()],
            ↪ s->length());
        } else if (dst_pos < 0) {
            ss.print("arraycopy: destination index %d out of bounds
            ↪ for %s[%d]", dst_pos,
            ↪ type2name_tab[ArrayKlass::cast(d->klass())->element_type()],
            ↪ d->length());
        } else {
            ss.print("arraycopy: length %d is negative", length);
        }

        ↪ THROW_MSG(vmSymbols::java_lang_ArrayIndexOutOfBoundsException(),
        ↪ ss.as_string());
    }

    // Check if the ranges are valid
    if (((unsigned int) length + (unsigned int) src_pos) >
    ↪ (unsigned int) s->length()) ||
    (((unsigned int) length + (unsigned int) dst_pos) > (unsigned
    ↪ int) d->length())) {

        // Pass specific exception reason.
        ResourceMark rm(THREAD);
        stringstream ss;
        if (((unsigned int) length + (unsigned int) src_pos) >
        ↪ (unsigned int) s->length()) {
            ss.print("arraycopy: last source index %u out of bounds
            ↪ for %s[%d]", (unsigned int) length + (unsigned int)
            ↪ src_pos,
            ↪ type2name_tab[ArrayKlass::cast(s->klass())->element_type()],
            ↪ s->length());
        } else {
            ss.print("arraycopy: last destination index %u out of
            ↪ bounds for %s[%d]", (unsigned int) length + (unsigned
            ↪ int) dst_pos,
            ↪ type2name_tab[ArrayKlass::cast(d->klass())->element_type()],
            ↪ d->length());
        }
    }

```

```

    ↪ THROW_MSG(vmSymbols::java_lang_ArrayIndexOutOfBoundsException(),
    ↪ ss.as_string());
}

// Check zero copy
if (length == 0)
    return;

// This is an attempt to make the copy_array fast.
int l2es = log2_element_size();
size_t src_offset =
    ↪ arrayOopDesc::base_offset_in_bytes(element_type()) +
    ↪ ((size_t)src_pos << l2es);
size_t dst_offset =
    ↪ arrayOopDesc::base_offset_in_bytes(element_type()) +
    ↪ ((size_t)dst_pos << l2es);
ArrayAccess<ARRAYCOPY_ATOMIC>::arraycopy<void>(s, src_offset,
    ↪ d, dst_offset, (size_t)length << l2es);
}

```

- hotspot/share/oops/access.hpp

```

// Helper for array access.
template <DecoratorSet decorators = INTERNAL_EMPTY>
class ArrayAccess: public HeapAccess<IS_ARRAY | decorators> {
    typedef HeapAccess<IS_ARRAY | decorators> AccessT;
public:
    template <typename T>
    static inline void arraycopy(arrayOop src_obj, size_t
    ↪ src_offset_in_bytes, arrayOop dst_obj, size_t
    ↪ dst_offset_in_bytes, size_t length) {
        AccessT::arraycopy(src_obj, src_offset_in_bytes,
        ↪ reinterpret_cast<const T*>(NULL), dst_obj,
        ↪ dst_offset_in_bytes, reinterpret_cast<T*>(NULL),
        ↪ length);
    }

    template <typename T>
    static inline void arraycopy_to_native(arrayOop src_obj,
    ↪ size_t src_offset_in_bytes, T* dst, size_t length) {
        AccessT::arraycopy(src_obj, src_offset_in_bytes,
        ↪ reinterpret_cast<const T*>(NULL), NULL, 0, dst,
        ↪ length);
    }

    template <typename T>

```

```

static inline void arraycopy_from_native(const T* src,
    ↪ arrayObj dst_obj, size_t dst_offset_in_bytes, size_t
    ↪ length) {
    AccessT::arraycopy(NULL, 0, src, dst_obj,
        ↪ dst_offset_in_bytes, reinterpret_cast<T*>(NULL),
        ↪ length);
}

static inline bool oop_arraycopy(arrayObj src_obj, size_t
    ↪ src_offset_in_bytes, arrayObj dst_obj, size_t
    ↪ dst_offset_in_bytes, size_t length) {
    return AccessT::oop_arraycopy(src_obj, src_offset_in_bytes,
        ↪ reinterpret_cast<const HeapWord*>(NULL), dst_obj,
        ↪ dst_offset_in_bytes, reinterpret_cast<HeapWord*>(NULL),
        ↪ length);
}

template <typename T>
static inline bool oop_arraycopy_raw(T* src, T* dst, size_t
    ↪ length) {
    return AccessT::oop_arraycopy(NULL, 0, src, NULL, 0, dst,
        ↪ length);
}

};

```

- hotspot/share/oops/access.hpp

```

template <DecoratorSet decorators = INTERNAL_EMPTY>
class Access: public AllStatic {
    // ...

protected:
    template <typename T>
    static inline bool oop_arraycopy(arrayObj src_obj, size_t
        ↪ src_offset_in_bytes, const T* src_raw, arrayObj dst_obj,
        ↪ size_t dst_offset_in_bytes, T* dst_raw, size_t length) {
        verify_decorators<ARRAYCOPY_DECORATOR_MASK | IN_HEAP |
            ↪ AS_DECORATOR_MASK | IS_ARRAY |
            ↪ IS_DEST_UNINITIALIZED>();
        return AccessInternal::arraycopy<decorators |
            ↪ INTERNAL_VALUE_IS_OOP>(src_obj, src_offset_in_bytes,
            ↪ src_raw, dst_obj, dst_offset_in_bytes, dst_raw,
            ↪ length);
    }

    template <typename T>

```

```

static inline void arraycopy(arrayOop src_obj, size_t
    ↪ src_offset_in_bytes, const T* src_raw, arrayOop dst_obj,
    ↪ size_t dst_offset_in_bytes, T* dst_raw, size_t length) {
    verify_decorators<ARRAYCOPY_DECORATOR_MASK | IN_HEAP |
        ↪ AS_DECORATOR_MASK | IS_ARRAY>();
    AccessInternal::arraycopy<decorators>(src_obj,
        ↪ src_offset_in_bytes, src_raw, dst_obj,
        ↪ dst_offset_in_bytes, dst_raw, length);
}
}

```

- hotspot/share/oops/accessBackend.inline.hpp

```

class RawAccessBarrierArrayCopy: public AllStatic {
    template<typename T> struct IsHeapWordSized: public
        ↪ IntegralConstant<bool, sizeof(T) == HeapWordSize> { };
public:
    template<DecoratorSet decorators, typename T>
    static inline typename EnableIf<
        HasDecorator<decorators,
            ↪ INTERNAL_VALUE_IS_OOP>::value>::type
    arraycopy(arrayOop src_obj, size_t src_offset_in_bytes, T*
        ↪ src_raw,
            arrayOop dst_obj, size_t dst_offset_in_bytes, T*
                ↪ dst_raw,
                    size_t length) {
        src_raw = arrayOopDesc::obj_offset_to_raw(src_obj,
            ↪ src_offset_in_bytes, src_raw);
        dst_raw = arrayOopDesc::obj_offset_to_raw(dst_obj,
            ↪ dst_offset_in_bytes, dst_raw);

        // We do not check for ARRAYCOPY_ATOMIC for oops, because
        ↪ they are unconditionally always atomic.
        if (HasDecorator<decorators, ARRAYCOPY_ARRAYOF>::value) {
            AccessInternal::arraycopy_arrayof_conjoint_oops(src_raw,
                ↪ dst_raw, length);
        } else {
            typedef typename HeapOopType<decorators>::type OopType;

            ↪ AccessInternal::arraycopy_conjoint_oops(reinterpret_cast<OopType*>(src_raw),
            ↪ reinterpret_cast<OopType*>(dst_raw), length);
        }
    }
}

template<DecoratorSet decorators, typename T>
static inline typename EnableIf<
    !HasDecorator<decorators, INTERNAL_VALUE_IS_OOP>::value &&
    HasDecorator<decorators, ARRAYCOPY_ARRAYOF>::value>::type
arraycopy(arrayOop src_obj, size_t src_offset_in_bytes, T*
    ↪ src_raw,

```

```

        arrayOp dst_obj, size_t dst_offset_in_bytes, T*
        ↪ dst_raw,
        size_t length) {
    src_raw = arrayOpDesc::obj_offset_to_raw(src_obj,
        ↪ src_offset_in_bytes, src_raw);
    dst_raw = arrayOpDesc::obj_offset_to_raw(dst_obj,
        ↪ dst_offset_in_bytes, dst_raw);

    AccessInternal::arraycopy_arrayof_conjoint(src_raw,
        ↪ dst_raw, length);
}

template <DecoratorSet decorators, typename T>
static inline typename EnableIf<
    !HasDecorator<decorators, INTERNAL_VALUE_IS_OOP>::value &&
    HasDecorator<decorators, ARRAYCOPY_DISJOINT>::value &&
    ↪ IsHeapWordSized<T>::value>::type
arraycopy(arrayOp src_obj, size_t src_offset_in_bytes, T*
    ↪ src_raw,
        arrayOp dst_obj, size_t dst_offset_in_bytes, T*
        ↪ dst_raw,
        size_t length) {
    src_raw = arrayOpDesc::obj_offset_to_raw(src_obj,
        ↪ src_offset_in_bytes, src_raw);
    dst_raw = arrayOpDesc::obj_offset_to_raw(dst_obj,
        ↪ dst_offset_in_bytes, dst_raw);

    // There is only a disjoint optimization for word
    ↪ granularity copying
    if (HasDecorator<decorators, ARRAYCOPY_ATOMIC>::value) {
        AccessInternal::arraycopy_disjoint_words_atomic(src_raw,
            ↪ dst_raw, length);
    } else {
        AccessInternal::arraycopy_disjoint_words(src_raw,
            ↪ dst_raw, length);
    }
}

template <DecoratorSet decorators, typename T>
static inline typename EnableIf<
    !HasDecorator<decorators, INTERNAL_VALUE_IS_OOP>::value &&
    !(HasDecorator<decorators, ARRAYCOPY_DISJOINT>::value &&
    ↪ IsHeapWordSized<T>::value) &&
    !HasDecorator<decorators, ARRAYCOPY_ARRAYOF>::value &&
    !HasDecorator<decorators, ARRAYCOPY_ATOMIC>::value>::type
arraycopy(arrayOp src_obj, size_t src_offset_in_bytes, T*
    ↪ src_raw,
        arrayOp dst_obj, size_t dst_offset_in_bytes, T*
        ↪ dst_raw,
        size_t length) {

```



```

src_raw = arrayOpDesc::obj_offset_to_raw(src_obj,
↳ src_offset_in_bytes, src_raw);
dst_raw = arrayOpDesc::obj_offset_to_raw(dst_obj,
↳ dst_offset_in_bytes, dst_raw);

AccessInternal::arraycopy_conjoint(src_raw, dst_raw,
↳ length);
}

template <DecoratorSet decorators, typename T>
static inline typename EnableIf<
!HasDecorator<decorators, INTERNAL_VALUE_IS_OOP>::value &&
!(HasDecorator<decorators, ARRAYCOPY_DISJOINT>::value &&
↳ IsHeapWordSized<T>::value) &&
!HasDecorator<decorators, ARRAYCOPY_ARRAYOF>::value &&
HasDecorator<decorators, ARRAYCOPY_ATOMIC>::value>::type
arraycopy(arrayOp src_obj, size_t src_offset_in_bytes, T*
↳ src_raw,
arrayOp dst_obj, size_t dst_offset_in_bytes, T*
↳ dst_raw,
size_t length) {
src_raw = arrayOpDesc::obj_offset_to_raw(src_obj,
↳ src_offset_in_bytes, src_raw);
dst_raw = arrayOpDesc::obj_offset_to_raw(dst_obj,
↳ dst_offset_in_bytes, dst_raw);

AccessInternal::arraycopy_conjoint_atomic(src_raw, dst_raw,
↳ length);
}
};

```

- hotspot/share/oops/accessBackend.cpp

```

// These forward copying calls to Copy without exposing the
↳ Copy type in headers unnecessarily

```

```

void arraycopy_arrayof_conjoint_oops(void* src, void* dst,
↳ size_t length) {
Copy::arrayof_conjoint_oops(reinterpret_cast<HeapWord*>(src),
↳ reinterpret_cast<HeapWord*>(dst), length);
}

```

```

void arraycopy_conjoint_oops(oop* src, oop* dst, size_t length)
↳ {
Copy::conjoint_oops_atomic(src, dst, length);
}

```

```

void arraycopy_conjoint_oops(narrowOop* src, narrowOop* dst,
↳ size_t length) {
Copy::conjoint_oops_atomic(src, dst, length);
}

```

```

}

void arraycopy_disjoint_words(void* src, void* dst, size_t
↪ length) {
    Copy::disjoint_words(reinterpret_cast<HeapWord*>(src),
↪ reinterpret_cast<HeapWord*>(dst), length);
}

void arraycopy_disjoint_words_atomic(void* src, void* dst,
↪ size_t length) {
    Copy::disjoint_words_atomic(reinterpret_cast<HeapWord*>(src),
↪ reinterpret_cast<HeapWord*>(dst), length);
}

template<>
void arraycopy_conjoint<jboolean>(jboolean* src, jboolean* dst,
↪ size_t length) {
    Copy::conjoint_jbytes(reinterpret_cast<jbyte*>(src),
↪ reinterpret_cast<jbyte*>(dst), length);
}

template<>
void arraycopy_conjoint<jbyte>(jbyte* src, jbyte* dst, size_t
↪ length) {
    Copy::conjoint_jbytes(src, dst, length);
}

template<>
void arraycopy_conjoint<jchar>(jchar* src, jchar* dst, size_t
↪ length) {
    Copy::conjoint_jshorts_atomic(reinterpret_cast<jshort*>(src),
↪ reinterpret_cast<jshort*>(dst), length);
}

template<>
void arraycopy_conjoint<jshort>(jshort* src, jshort* dst,
↪ size_t length) {
    Copy::conjoint_jshorts_atomic(src, dst, length);
}

template<>
void arraycopy_conjoint<jint>(jint* src, jint* dst, size_t
↪ length) {
    Copy::conjoint_jints_atomic(src, dst, length);
}

template<>
void arraycopy_conjoint<jfloat>(jfloat* src, jfloat* dst,
↪ size_t length) {

```

```

    Copy::conjoint_jints_atomic(reinterpret_cast<jint*>(src),
    ↪ reinterpret_cast<jint*>(dst), length);
}

template<>
void arraycopy_conjoint<jlong>(jlong* src, jlong* dst, size_t
    ↪ length) {
    Copy::conjoint_jlongs_atomic(src, dst, length);
}

template<>
void arraycopy_conjoint<jdouble>(jdouble* src, jdouble* dst,
    ↪ size_t length) {
    Copy::conjoint_jlongs_atomic(reinterpret_cast<jlong*>(src),
    ↪ reinterpret_cast<jlong*>(dst), length);
}

template<>
void arraycopy_arrayof_conjoint<jbyte>(jbyte* src, jbyte* dst,
    ↪ size_t length) {

    ↪ Copy::arrayof_conjoint_jbytes(reinterpret_cast<HeapWord*>(src),
    ↪ reinterpret_cast<HeapWord*>(dst), length);
}

template<>
void arraycopy_arrayof_conjoint<jshort>(jshort* src, jshort*
    ↪ dst, size_t length) {

    ↪ Copy::arrayof_conjoint_jshorts(reinterpret_cast<HeapWord*>(src),
    ↪ reinterpret_cast<HeapWord*>(dst),
    ↪ length);
}

template<>
void arraycopy_arrayof_conjoint<jint>(jint* src, jint* dst,
    ↪ size_t length) {

    ↪ Copy::arrayof_conjoint_jints(reinterpret_cast<HeapWord*>(src),
    ↪ reinterpret_cast<HeapWord*>(dst), length);
}

template<>
void arraycopy_arrayof_conjoint<jlong>(jlong* src, jlong* dst,
    ↪ size_t length) {

    ↪ Copy::arrayof_conjoint_jlongs(reinterpret_cast<HeapWord*>(src),
    ↪ reinterpret_cast<HeapWord*>(dst), length);
}

```

```

template<>
void arraycopy_conjoint<void>(void* src, void* dst, size_t
↪ length) {
    Copy::conjoint_jbytes(reinterpret_cast<jbyte*>(src),
↪ reinterpret_cast<jbyte*>(dst), length);
}

template<>
void arraycopy_conjoint_atomic<jbyte>(jbyte* src, jbyte* dst,
↪ size_t length) {
    Copy::conjoint_jbytes_atomic(src, dst, length);
}

template<>
void arraycopy_conjoint_atomic<jshort>(jshort* src, jshort*
↪ dst, size_t length) {
    Copy::conjoint_jshorts_atomic(src, dst, length);
}

template<>
void arraycopy_conjoint_atomic<jint>(jint* src, jint* dst,
↪ size_t length) {
    Copy::conjoint_jints_atomic(src, dst, length);
}

template<>
void arraycopy_conjoint_atomic<jlong>(jlong* src, jlong* dst,
↪ size_t length) {
    Copy::conjoint_jlongs_atomic(src, dst, length);
}

template<>
void arraycopy_conjoint_atomic<void>(void* src, void* dst,
↪ size_t length) {
    Copy::conjoint_memory_atomic(src, dst, length);
}

```

- hotspot/share/utilities/copy.hpp

```

class Copy : AllStatic {
public:
    // Block copy methods have four attributes. We don't define
    ↪ all possibilities.
    // alignment: aligned to BytesPerLong
    // arrayof: arraycopy operation with both operands
    ↪ aligned on the same
    // boundary as the first element of an array of
    ↪ the copy unit.
    // This is currently a HeapWord boundary on all
    ↪ platforms, except

```

```

//          for long and double arrays, which are aligned
→ on an 8-byte
//          boundary on all platforms.
//          arraycopy operations are implicitly atomic on
→ each array element.
// overlap: disjoint or conjoint.
// copy unit: bytes or words (i.e., HeapWords) or oops
→ (i.e., pointers).
// atomicity: atomic or non-atomic on the copy unit.
//
// Names are constructed thusly:
//
//      [ 'aligned_' | 'arrayof_' ]
//      ('conjoint_' | 'disjoint_')
//      ('words' | 'bytes' | 'jshorts' | 'jints' | 'jlongs' |
→ 'oops')
//      [ '_atomic' ]
//
// Except in the arrayof case, whatever the alignment is, we
→ assume we can copy
// whole alignment units. E.g., if BytesPerLong is 2x word
→ alignment, an odd
// count may copy an extra word. In the arrayof case, we are
→ allowed to copy
// only the number of copy units specified.
//
// All callees check count for 0.
//

// HeapWords

// Word-aligned words,    conjoint, not atomic on each word
static void conjoint_words(const HeapWord* from, HeapWord*
→ to, size_t count) {
    assert_params_ok(from, to, HeapWordSize);
    pd_conjoint_words(from, to, count);
}

// Word-aligned words,    disjoint, not atomic on each word
static void disjoint_words(const HeapWord* from, HeapWord*
→ to, size_t count) {
    assert_params_ok(from, to, HeapWordSize);
    assert_disjoint(from, to, count);
    pd_disjoint_words(from, to, count);
}

// Word-aligned words,    disjoint, atomic on each word
static void disjoint_words_atomic(const HeapWord* from,
→ HeapWord* to, size_t count) {
    assert_params_ok(from, to, HeapWordSize);

```

```

    assert_disjoint(from, to, count);
    pd_disjoint_words_atomic(from, to, count);
}

// Object-aligned words, conjoint, not atomic on each word
static void aligned_conjoint_words(const HeapWord* from,
    ↪ HeapWord* to, size_t count) {
    assert_params_aligned(from, to);
    pd_aligned_conjoint_words(from, to, count);
}

// Object-aligned words, disjoint, not atomic on each word
static void aligned_disjoint_words(const HeapWord* from,
    ↪ HeapWord* to, size_t count) {
    assert_params_aligned(from, to);
    assert_disjoint(from, to, count);
    pd_aligned_disjoint_words(from, to, count);
}

// bytes, jshorts, jint, jlongs, oops

// bytes,                conjoint, not atomic on each byte
    ↪ (not that it matters)
static void conjoint_jbytes(const void* from, void* to,
    ↪ size_t count) {
    pd_conjoint_bytes(from, to, count);
}

// bytes,                conjoint, atomic on each byte (not
    ↪ that it matters)
static void conjoint_jbytes_atomic(const void* from, void*
    ↪ to, size_t count) {
    pd_conjoint_bytes(from, to, count);
}

// jshorts,              conjoint, atomic on each jshort
static void conjoint_jshorts_atomic(const jshort* from,
    ↪ jshort* to, size_t count) {
    assert_params_ok(from, to, BytesPerShort);
    pd_conjoint_jshorts_atomic(from, to, count);
}

// jint,                 conjoint, atomic on each jint
static void conjoint_jints_atomic(const jint* from, jint* to,
    ↪ size_t count) {
    assert_params_ok(from, to, BytesPerInt);
    pd_conjoint_jints_atomic(from, to, count);
}

// jlongs,              conjoint, atomic on each jlong

```

```

static void conjoint_jlongs_atomic(const jlong* from, jlong*
→ to, size_t count) {
    assert_params_ok(from, to, BytesPerLong);
    pd_conjoint_jlongs_atomic(from, to, count);
}

// oops,                conjoint, atomic on each oop
static void conjoint_oops_atomic(const oop* from, oop* to,
→ size_t count) {
    assert_params_ok(from, to, BytesPerHeapOop);
    pd_conjoint_oops_atomic(from, to, count);
}

// overloaded for UseCompressedOops
static void conjoint_oops_atomic(const narrowOop* from,
→ narrowOop* to, size_t count) {
    assert(sizeof(narrowOop) == sizeof(jint), "this cast is
→ wrong");
    assert_params_ok(from, to, BytesPerInt);
    pd_conjoint_jints_atomic((const jint*)from, (jint*)to,
→ count);
}

// Copy a span of memory.  If the span is an integral number
→ of aligned
// longs, words, or ints, copy those units atomically.
// The largest atomic transfer unit is 8 bytes, or the
→ largest power
// of two which divides all of from, to, and size, whichever
→ is smaller.
static void conjoint_memory_atomic(const void* from, void*
→ to, size_t size);

// bytes,                conjoint array, atomic on each byte
→ (not that it matters)
static void arrayof_conjoint_jbytes(const HeapWord* from,
→ HeapWord* to, size_t count) {
    pd_arrayof_conjoint_bytes(from, to, count);
}

// jshorts,             conjoint array, atomic on each
→ jshort
static void arrayof_conjoint_jshorts(const HeapWord* from,
→ HeapWord* to, size_t count) {
    assert_params_ok(from, to, BytesPerShort);
    pd_arrayof_conjoint_jshorts(from, to, count);
}

// jints,                conjoint array, atomic on each jint

```

```

static void arrayof_conjoint_jints(const HeapWord* from,
    ↪ HeapWord* to, size_t count) {
    assert_params_ok(from, to, BytesPerInt);
    pd_arrayof_conjoint_jints(from, to, count);
}

// jlongs,                conjoint array, atomic on each
    ↪ jlong
static void arrayof_conjoint_jlongs(const HeapWord* from,
    ↪ HeapWord* to, size_t count) {
    assert_params_ok(from, to, BytesPerLong);
    pd_arrayof_conjoint_jlongs(from, to, count);
}

// oops,                conjoint array, atomic on each oop
static void arrayof_conjoint_oops(const HeapWord* from,
    ↪ HeapWord* to, size_t count) {
    assert_params_ok(from, to, BytesPerHeapOop);
    pd_arrayof_conjoint_oops(from, to, count);
}

// Known overlap methods

// Copy word-aligned words from higher to lower addresses,
    ↪ not atomic on each word
inline static void conjoint_words_to_lower(const HeapWord*
    ↪ from, HeapWord* to, size_t byte_count) {
    // byte_count is in bytes to check its alignment
    assert_params_ok(from, to, HeapWordSize);
    assert_byte_count_ok(byte_count, HeapWordSize);

    size_t count = align_up(byte_count, HeapWordSize) >>
    ↪ LogHeapWordSize;
    assert(to <= from || from + count <= to, "do not overwrite
    ↪ source data");

    while (count-- > 0) {
        *to++ = *from++;
    }
}

// Copy word-aligned words from lower to higher addresses,
    ↪ not atomic on each word
inline static void conjoint_words_to_higher(const HeapWord*
    ↪ from, HeapWord* to, size_t byte_count) {
    // byte_count is in bytes to check its alignment
    assert_params_ok(from, to, HeapWordSize);
    assert_byte_count_ok(byte_count, HeapWordSize);

```



```

size_t count = align_up(byte_count, HeapWordSize) >>
↳ LogHeapWordSize;
assert(from <= to || to + count <= from, "do not overwrite
↳ source data");

from += count - 1;
to   += count - 1;
while (count-- > 0) {
    *to-- = *from--;
}
}

/**
 * Copy elements
 *
 * @param src address of source
 * @param dst address of destination
 * @param byte_count number of bytes to copy
 * @param elem_size size of the elements to copy-swap
 */
static void conjoint_copy(const void* src, void* dst, size_t
↳ byte_count, size_t elem_size);

/**
 * Copy and *unconditionally* byte swap elements
 *
 * @param src address of source
 * @param dst address of destination
 * @param byte_count number of bytes to copy
 * @param elem_size size of the elements to copy-swap
 */
static void conjoint_swap(const void* src, void* dst, size_t
↳ byte_count, size_t elem_size);

/**
 * Copy and byte swap elements from the specified endian to
↳ the native (cpu) endian if needed (if they differ)
 *
 * @param src address of source
 * @param dst address of destination
 * @param byte_count number of bytes to copy
 * @param elem_size size of the elements to copy-swap
 */
template <Endian::Order endian>
static void conjoint_swap_if_needed(const void* src, void*
↳ dst, size_t byte_count, size_t elem_size) {
    if (Endian::NATIVE != endian) {
        conjoint_swap(src, dst, byte_count, elem_size);
    } else {
        conjoint_copy(src, dst, byte_count, elem_size);
    }
}

```

```

    }
}

// Fill methods

// Fill word-aligned words, not atomic on each word
// set_words
static void fill_to_words(HeapWord* to, size_t count, jint
    ↪ value = 0) {
    assert_params_ok(to, HeapWordSize);
    pd_fill_to_words(to, count, value);
}

static void fill_to_aligned_words(HeapWord* to, size_t count,
    ↪ jint value = 0) {
    assert_params_aligned(to);
    pd_fill_to_aligned_words(to, count, value);
}

// Fill bytes
static void fill_to_bytes(void* to, size_t count, jubyte
    ↪ value = 0) {
    pd_fill_to_bytes(to, count, value);
}

// Fill a span of memory. If the span is an integral number
    ↪ of aligned
// longs, words, or ints, store to those units atomically.
// The largest atomic transfer unit is 8 bytes, or the
    ↪ largest power
// of two which divides both to and size, whichever is
    ↪ smaller.
static void fill_to_memory_atomic(void* to, size_t size,
    ↪ jubyte value = 0);

// Zero-fill methods

// Zero word-aligned words, not atomic on each word
static void zero_to_words(HeapWord* to, size_t count) {
    assert_params_ok(to, HeapWordSize);
    pd_zero_to_words(to, count);
}

// Zero bytes
static void zero_to_bytes(void* to, size_t count) {
    pd_zero_to_bytes(to, count);
}

private:

```

```

static bool params_disjoint(const HeapWord* from, HeapWord*
↪ to, size_t count) {
    if (from < to) {
        return pointer_delta(to, from) >= count;
    }
    return pointer_delta(from, to) >= count;
}

// These methods raise a fatal if they detect a problem.

static void assert_disjoint(const HeapWord* from, HeapWord*
↪ to, size_t count) {
    assert(params_disjoint(from, to, count), "source and dest
↪ overlap");
}

static void assert_params_ok(const void* from, void* to,
↪ intptr_t alignment) {
    assert(is_aligned(from, alignment), "must be aligned: "
↪ INTPTR_FORMAT, p2i(from));
    assert(is_aligned(to, alignment), "must be aligned: "
↪ INTPTR_FORMAT, p2i(to));
}

static void assert_params_ok(HeapWord* to, intptr_t
↪ alignment) {
    assert(is_aligned(to, alignment), "must be aligned: "
↪ INTPTR_FORMAT, p2i(to));
}

static void assert_params_aligned(const HeapWord* from,
↪ HeapWord* to) {
    assert(is_aligned(from, BytesPerLong), "must be aligned: "
↪ INTPTR_FORMAT, p2i(from));
    assert(is_aligned(to, BytesPerLong), "must be aligned: "
↪ INTPTR_FORMAT, p2i(to));
}

static void assert_params_aligned(HeapWord* to) {
    assert(is_aligned(to, BytesPerLong), "must be aligned: "
↪ INTPTR_FORMAT, p2i(to));
}

static void assert_byte_count_ok(size_t byte_count, size_t
↪ unit_size) {
    assert(is_aligned(byte_count, unit_size), "byte count must
↪ be aligned");
}

// Platform dependent implementations of the above methods.

```

```

#include CPU_HEADER(copy)

};

• hotspot/cpu/ppc/copy_ppc.hpp

// Inline functions for memory copy and fill.


static void pd_conjoint_words(const HeapWord* from, HeapWord*
↪ to, size_t count) {
    (void)memmove(to, from, count * HeapWordSize);
}

// Template for atomic, element-wise copy.
template <class T>
static void copy_conjoint_atomic(const T* from, T* to, size_t
↪ count) {
    if (from > to) {
        while (count-- > 0) {
            // Copy forwards
            *to++ = *from++;
        }
    } else {
        from += count - 1;
        to += count - 1;
        while (count-- > 0) {
            // Copy backwards
            *to-- = *from--;
        }
    }
}

// ...


```

0.3.1.6 Type Array Inheritance




0.3.1.7 TLAB

0.3.1.7.1 Inheritance



0.3.1.7.2 Initialization



1. `void OptoRuntime::new_array_C(Klass* array_type, int len, JavaThread *thread)`
2. `typeArrayOop oopFactory::new_typeArray(BasicType type, int length, TRAPS)`
3. `TypeArrayKlass* TypeArrayKlass::allocate(ClassLoaderData* loader_data, BasicType type, Symbol* name, TRAPS)`
4. `TypeArrayKlass* TypeArrayKlass::allocate(ClassLoaderData* loader_data, BasicType type, Symbol* name, TRAPS)`
5. `objArrayOop oopFactory::new_objArray(Klass* klass, int length, TRAPS)`

0.3.1.7.3 Notes

The virtual machine must never call one of the implicitly declared global allocation or deletion functions. (Such calls may result in link-time or run-time errors.) For convenience and documentation of intended use, classes in the virtual machine may be derived from one of the following allocation classes, some of which define allocation and deletion functions. Note: `std::malloc` and `std::free` should never be called directly. For objects allocated in the resource area (see `resourceArea.hpp`). - `ResourceObj`

For objects allocated in the C-heap (managed by: free & malloc and tracked with NMT) - CHeapObj

For objects allocated on the stack. - StackObj

For classes used as name spaces. - AllStatic

For classes in Metaspace (class data) - MetaspaceObj

The printable subclasses are used for debugging and define virtual member functions for printing. Classes that avoid allocating the vtbl entries in the objects should therefore not be the printable subclasses.

The following macros and function should be used to allocate memory directly in the resource area or in the C-heap, The _OBJ variants of the NEW/FREE_C_HEAP macros are used for alloc/dealloc simple objects which are not inherited from CHeapObj, note constructor and destructor are not called. The preferable way to allocate objects is using the new operator.

WARNING: The array variant must only be used for a homogenous array where all objects are of the exact type specified. If subtypes are stored in the array then must pay attention to calling destructors at needed.

- NEW_RESOURCE_ARRAY(type, size)
- NEW_RESOURCE_OBJ(type)
- NEW_C_HEAP_ARRAY(type, size)
- NEW_C_HEAP_OBJ(type, memflags)
- FREE_C_HEAP_ARRAY(type, old)
- FREE_C_HEAP_OBJ(objname, type, memflags)
- `char* AllocateHeap(size_t size, const char* name);`
- `void FreeHeap(void* p);`

In non product mode we introduce a super class for all allocation classes that supports printing. We avoid the superclass in product mode to save space.

0.4 Notes

0.4.1 Tuesday, 08 October 2019

0.4.1.1 Kafka Bug Study

- Distributed Concurrency Bugs
 1. [\[KAFKA-1183\]](#) - `DefaultEventHandler` causes unbalanced distribution of messages across partitions
 2. [\[KAFKA-1154\]](#) - replicas may not have consistent data after becoming follower
 3. [\[KAFKA-1124\]](#) - Sending to a new topic (with `auto.create.topics.enable`) returns `ERROR`
 4. [\[KAFKA-1097\]](#) - Race condition while reassigning low throughput partition leads to incorrect ISR information in zookeeper
- Scalability Bugs
 1. [\[KAFKA-1228\]](#) - Socket Leak on `ReplicaFetcherThread`

0.4.1.2 JVM Deduplication

- Trying to implement the `madvise` call inside the TLAB allocation (`ThreadLocalAllocBuffer::allocate`)
Only saving 44kb for two JVM process and array 139 with its elements 42

0.4.2 Friday, 04 October 2019

0.4.2.1 Allocation Classes

1. Allocation, consists of:
 - 1.1. AllocatedObj (Abstract Class)
 - 1.1.1. CHeapObj (Abstract Class)
 - 1.1.1.1. CollectedHeap
 - 1.1.1.1.1. GenCollectedHeap
 - 1.1.1.1.1.1. SerialHeap
 - 1.1.1.1.1.2. CMSHeap
 - 1.1.1.1.1.2. G1CollectedHeap
 - 1.1.1.1.3. ParallelScavengeHeap
 - 1.1.1.1.4. ZCollectedHeap
 - 1.1.2. StackObj (Abstract Class)
 - 1.1.2.1. MemAllocator
 - 1.1.2.1.1. ObjAllocator
 - 1.1.2.1.2. ObjArrayAllocator
 - 1.1.2.1.3. ClassAllocator
 - 1.1.3. ResourceObj (Abstract Class)
 - 1.2. MetaspaceObj (Abstract Class)
 - 1.3. AllStatic (Abstract Class)
 - 1.3.1. ArrayAllocator
 - 1.3.2. MmapArrayAllocator
 - 1.4. MallocArrayAllocator

0.4.2.2 Notes

- All of the Klass is being instantiated in Metaspace
- EpsilonGC is not creating its own heap

0.4.3 Saturday, 19 October 2019

0.4.3.1 JVM Deduplication

- Appending this code into the ‘accessBackend.inline.hpp’ especially in function ‘arraycopy’ but for atomic one

```
// @rayandrew
// added this to add execute ksm

if (os::can_execute_ksm() && check_if_tescase_array(length)) {
    // tty->print_cr("Pointer actual size " SIZE_FORMAT "Type
    → %s", sizeof(*dst_raw), demangle(typeid(dst_raw).name()));
    long page_size = sysconf(_SC_PAGE_SIZE);
    tty->print_cr("Is aligned : %d",
    → is_aligned(dst_obj->base(T_BYTE), page_size));
    tty->print_cr("Raw pointer : " PTR_FORMAT ", DST_OBJ "
    → PTR_FORMAT, dst_raw, dst_obj->base(T_BYTE));
    // tty->print_cr("Pointer actual size " SIZE_FORMAT "Type
    → %s", sizeof(*dst_raw), demangle(typeid(dst_raw).name()));
    tty->print_cr("Value : %s", dst_obj->print_value_string());
    tty->print_cr("Raw pointer : " PTR_FORMAT ", DST_OBJ "
    → PTR_FORMAT, dst_raw, dst_obj->base(T_BYTE));
    tty->print_cr("Length : %d", dst_obj->length());
    tty->print_cr("Sizeof arrayOopDesc : %d",
    → sizeof(arrayOopDesc));
    // tty->print_cr("Content %d", &dst_raw[0]);
    tty->print_cr("Element offset %d",
    → (size_t)typeArrayOopDesc::element_offset<jbyte>(2));
    tty->print_cr("New Content 0 ptr " PTR_FORMAT " value %d",
    → static_cast<typeArrayOop>(dst_obj)->byte_at_addr(0),
    → (int)*static_cast<typeArrayOop>(dst_obj)->byte_at_addr(0));
    tty->print_cr("New Content 1 ptr " PTR_FORMAT " value %d",
    → static_cast<typeArrayOop>(dst_obj)->byte_at_addr(1),
    → (int)*static_cast<typeArrayOop>(dst_obj)->byte_at_addr(1));
    tty->print_cr("Content 0 %d", *(int*)((intptr_t)dst_obj));
    tty->print_cr("Content 1 %d", *(int*)((intptr_t)dst_obj +
    → 17));
    tty->print_cr("Content %d %d", dst_obj->length(),
    → *(int*)((intptr_t)dst_obj + dst_obj->length()));

    → os::mark_for_mergeable_debug(static_cast<typeArrayOop>(dst_obj)->byte_at_addr(0),
    → length, "RawAccessBarrierArrayCopy::arraycopy [5]");
    // os::mark_for_mergeable_debug(dst_obj->base(T_BYTE),
    → length, "RawAccessBarrierArrayCopy::arraycopy [5]");
}
```

- And the result is as follow :

```
KSM features is enabled
[0.010s][info][gc] Resizeable heap; starting at 2011M, max:
→ 30718M, step: 128M
```

```

[0.010s][info][gc] Using TLAB allocation; max: 4096K
[0.010s][info][gc] Elastic TLABs enabled; elasticity: 1.10x
[0.010s][info][gc] Elastic TLABs decay enabled; decay time:
↳ 1000ms
[0.010s][info][gc] Using Epsilon
Phase 1
Phase 2, Copy array
Is aligned : 0
Raw pointer : 0x0000000080697148, DST_OBJ 0x0000000080697148
Value : [B{0x0000000080697138}]
Raw pointer : 0x0000000080697148, DST_OBJ 0x0000000080697148
Length : 4096000
Sizeof arrayOpDesc : 16
Element offset 18
New Content 0 ptr 0x0000000080697148 value 1
New Content 1 ptr 0x0000000080697149 value 1
Content 0 1
Content 1 16843009
Content 4096000 16843009
[RawAccessBarrierArrayCopy::arraycopy [5]] Mark page address :
↳ 0x0000000080697148 bytes : 4096000 to be mergeable
OpenJDK 64-Bit Server VM warning: Failed to memory pages using
↳ madvise, req_addr: 0x0000000080697148 bytes: 4096000 (errno
↳ = 22).
Phase 3
src == dst : true
^C[5.057s][info][gc] Heap: 30718M reserved, 2011M (6.55%)
↳ committed, 9170K (0.03%) used

```

- The result that can we take Array is not page aligned : this is actually mandatory in MADV MERGEABLE

0.5 Resources

0.5.1 Kernel Samepage Merging

1. Documentation

- (a) [Madvise](#)
- (b) [KSM](#)
- (c) [Metaspace Architecture](#)